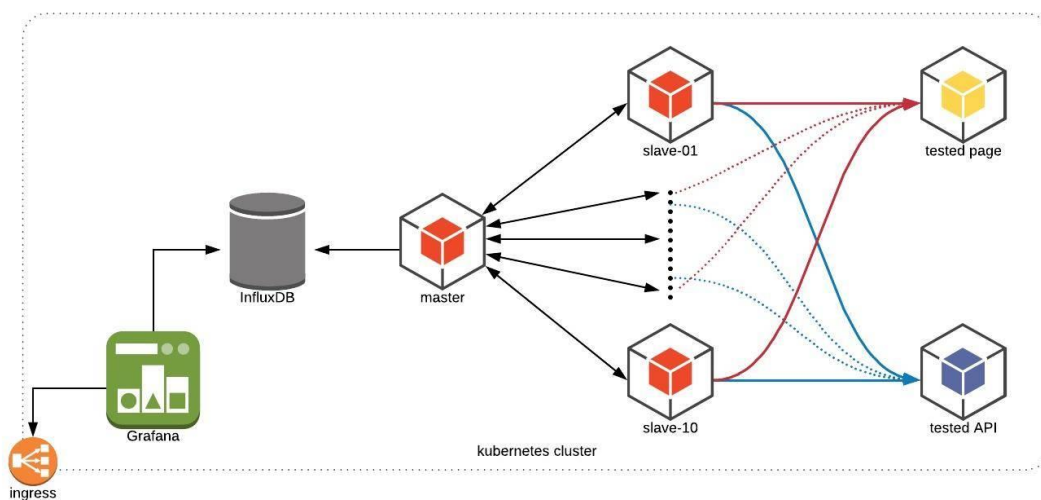


## 一、 前提条件

[Kubernetes](#) > 1.16

## 二、 部署拓扑



可以从 master 节点启动测试，master 节点把对应的测试脚本发送到对应的 slaves 节点，slave 节点的 pod/nodes 主要作用即发压。

部署文件清单：

- `jmeter_cluster_create.sh` — 此脚本将要求一个唯一的 namespace，然后它将继续创建命名空间和所有组件（jmeter master, slaves, influxdb 和 grafana）。
- 注意：在启动前，请在 `jmeter_slaves_deploy.yaml` 文件中设置要用于 slaves 服务器的副本数，通常副本数应与拥有的 worker nodes 相匹配。
- `jmeter_master_configmap.yaml` — Jmeter master 的应用配置。
- `jmeter_master_deployment.yaml` — Jmeter master 的部署清单。
- `jmeter_slaves_deploy.yaml` — Jmeter slave 的部署清单。
- `jmeter_slave_svc.yaml` — jmeter slave 的服务清单。使用 headless service，这使我们能够直接获取 jmeter slave 的 POD IP 地址，而我们不需要 DNS 或轮询。创建此文件是为了使 slave Pod IP 地址更容易直接发送到 jmeter master。
- `jmeter_influxdb_configmap.yaml` — influxdb 部署的应用配置。如果要在默认的内fluxdb 端口之外使用 graphite 存储方法，这会将 influxdb 配置为暴露端口 2003，以便支持 graphite。因此，可以使用 influxdb 部署来支持 jmeter 后置监听器方法（graphite 和 influxdb）。
- `jmeter_influxdb_deploy.yaml` — Influxdb 的部署清单
- `jmeter_influxdb_svc.yaml` — Influxdb 的服务清单。
- `jmeter_grafana_deploy.yaml` — grafana 部署清单。
- `jmeter_grafana_svc.yaml` — grafana 部署的服务清单，默认情况下使用 NodePort，如果公有云中运行它，则可以将其更改为 LoadBalancer（并且可以设置 CNAME 以使用 FQDN 缩短名称）。
- `dashboard.sh` — 该脚本用于自动创建以下内容：
  - (1) influxdb pod 中的一个 influxdb 数据库（Jmeter）
  - (2) grafana 中的数据源（Jmeterdb）
- `start_test.sh` — 此脚本用于自动运行 Jmeter 测试脚本，而无需手动登录 Jmeter 主 shell，它将询问 Jmeter 测试脚本的位置，然后将其复制到 Jmeter master pod 并启动自动对 Jmeter slave 进行测试。
- `jmeter_stop.sh` - 停止测试
- `GrafanaJMeterTemplate.json` — 预先构建的 Jmeter grafana 仪表盘。
- `Dockerfile-base` - 构建 Jmeter 基础镜像
- `Dockerfile-master` - 构建 Jmeter master 镜像
- `Dockerfile-slave` - 构建 Jmeter slave 镜像
- `Dockerimages.sh` - 批量构建 docker 镜像

## 三、docker 镜像

### 1、构建 docker 镜像

执行脚本，构建镜像：

```
1 | ./dockerimages.sh
```

查看镜像：

```
1 | $ docker images
```

将镜像推送到 Registry：

```
1 | $ sudo docker login --username=xxxx registry.cn-beijing.aliyuncs.com
2 | $ sudo docker tag [ImageId] registry.cn-beijing.aliyuncs.com/7d/jmeter-base:[镜像版本号]
3 | $ sudo docker push registry.cn-beijing.aliyuncs.com/7d/jmeter-base:[镜像版本号]
```

### 2、部署清单

Dockerfile-base （构建 [Jmeter](#) 基础镜像）：

```
1 | FROM alpine:latest
2 | LABEL MAINTAINER 7DGroup
3 |
4 | ARG JMETER_VERSION=5.2.1
5 |
6 | #定义时区参数
7 | ENV TZ=Asia/Shanghai
8 |
9 | RUN apk update && \
10 |    apk upgrade && \
11 |    apk add --update openjdk8-jre wget tar bash && \
12 |    mkdir /jmeter && cd /jmeter/ && \
13 |    wget https://mirrors.tuna.tsinghua.edu.cn/apache/jmeter/binaries/apache-jmeter-${JMETER_VERSION}.tgz &
14 |    tar -xzf apache-jmeter-${JMETER_VERSION}.tgz && rm apache-jmeter-${JMETER_VERSION}.tgz && \
15 |    cd /jmeter/apache-jmeter-${JMETER_VERSION}/ && \
16 |    wget -q -O /tmp/JMeterPlugins-Standard-1.4.0.zip https://jmeter-plugins.org/downloads/file/JMeterPlugi
17 |    wget -q -O /jmeter/apache-jmeter-${JMETER_VERSION}/lib/ext/pepper-box-1.0.jar https://github.com/raladev
18 |    cd /jmeter/apache-jmeter-${JMETER_VERSION}/ && \
19 |    wget -q -O /tmp/bzm-parallel-0.7.zip https://jmeter-plugins.org/files/packages/bzm-parallel-0.7.zip &&
20 |    ln -snf /usr/share/zoneinfo/$TZ /etc/localtime && echo "$TZ" > /etc/timezone
21 |
22 | ENV JMETER_HOME /jmeter/apache-jmeter-${JMETER_VERSION}/
23 |
24 | ENV PATH $JMETER_HOME/bin:$PATH
25 | #JMeter 主配置文件
26 | ADD jmeter.properties $JMETER_HOME/bin/jmeter.properties
```

Dockerfile-master (构建 Jmeter master 镜像) :

```
1 FROM registry.cn-beijing.aliyuncs.com/7d/jmeter-base:latest
2 MAINTAINER 7DGroup
3
4 EXPOSE 60000
```

Dockerfile-slave (构建 Jmeter slave 镜像) :

```
1 Dockerfile-slave:
2 FROM registry.cn-beijing.aliyuncs.com/7d/jmeter-base:latest
3 MAINTAINER 7DGroup
4
5 EXPOSE 1099 50000
6
7 ENTRYPOINT $JMETER_HOME/bin/jmeter-server \
8 -Dserver.rmi.localport=50000 \
9 -Dserver_port=1099 \
10 -Dserver.rmi.ssl.disable=true
```

Dockerimages.sh (批量构建 docker 镜像) :

```
1 #!/bin/bash -e
2 docker build --tag="registry.cn-beijing.aliyuncs.com/7d/jmeter-base:latest" -f Dockerfile-base .
3 docker build --tag="registry.cn-beijing.aliyuncs.com/7d/jmeter-master:latest" -f Dockerfile-master .
4 docker build --tag="registry.cn-beijing.aliyuncs.com/7d/jmeter-slave:latest" -f Dockerfile-slave .
```

## 四、Kubernetes 部署

### 1、部署组件

执行 `jmeter_cluster_create.sh`, 并输入一个唯一的 namespace

```
1 | ./jmeter_cluster_create.sh
```

等待一会, 查看pods安装情况:

```
1 $ kubectl get pods -n 7dgroup
2 NAME                                READY   STATUS    RESTARTS   AGE
3 influxdb-jmeter-584cf69759-j5m85   1/1     Running   2           5m
4 jmeter-grafana-6d5b75b7f6-57dxj     1/1     Running   1           5m
5 jmeter-master-84bfd5d96d-kthzm      1/1     Running   0           5m
6 jmeter-slaves-b5b75757-dxkxz        1/1     Running   0           5m
7 jmeter-slaves-b5b75757-n58jw        1/1     Running   0           5m
```

### 2、部署清单

## 2.1、主执行脚本

jmeter\_cluster\_create.sh (创建命名空间和所有组件 (jmeter master, slaves, influxdb 和 grafana) ) :

```
1  #!/usr/bin/env bash
2  #Create multiple Jmeter namespaces on an existing kuberntes cluster
3  #Started On January 23, 2018
4  working_dir=`pwd`
5  echo "checking if kubectl is present"
6  if ! hash kubectl 2>/dev/null
7  then
8  echo "'kubectl' was not found in PATH"
9  echo "Kindly ensure that you can acces an existing kubernetes cluster via kubectl"
10 exit
11 fi
12 kubectl version --short
13 echo "Current list of namespaces on the kubernetes cluster:"
14 echo
15 kubectl get namespaces | grep -v NAME | awk '{print $1}'
16 echo
17 echo "Enter the name of the new tenant unique name, this will be used to create the namespace"
18 read tenant
19 echo
20 #Check If namespace exists
21 kubectl get namespace $tenant > /dev/null 2>&1
22 if [ $? -eq 0 ]
23 then
24 echo "Namespace $tenant already exists, please select a unique name"
25 echo "Current list of namespaces on the kubernetes cluster"
26 sleep 2
27 kubectl get namespaces | grep -v NAME | awk '{print $1}'
28 exit 1
29 fi
30 echo
31 echo "Creating Namespace: $tenant"
32 kubectl create namespace $tenant
33 echo "Namspace $tenant has been created"
34 echo
35 echo "Creating Jmeter slave nodes"
36 nodes=`kubectl get no | egrep -v "master|NAME" | wc -l`
37 echo
38 echo "Number of worker nodes on this cluster is " $nodes
39 echo
40 #echo "Creating $nodes Jmeter slave replicas and service"
41 echo
42 kubectl create -n $tenant -f $working_dir/jmeter_slaves_deploy.yaml
43 kubectl create -n $tenant -f $working_dir/jmeter_slaves_svc.yaml
44 echo "Creating Jmeter Master"
45 kubectl create -n $tenant -f $working_dir/jmeter_master_configmap.yaml
46 kubectl create -n $tenant -f $working_dir/jmeter_master_deploy.yaml
47
48 echo "Creating Influxdb and the service"
49 kubectl create -n $tenant -f $working_dir/jmeter_influxdb_configmap.yaml
50 kubectl create -n $tenant -f $working_dir/jmeter_influxdb_deploy.yaml
51 kubectl create -n $tenant -f $working_dir/jmeter_influxdb_svc.yaml
52 echo "Creating Grafana Deployment"
53 kubectl create -n $tenant -f $working_dir/jmeter_grafana_deploy.yaml
54 kubectl create -n $tenant -f $working_dir/jmeter_grafana_svc.yaml
55 echo "Printout Of the $tenant Objects"
56 echo
57 kubectl get -n $tenant all
58 echo namespace = $tenant > $working_dir/tenant_export
```

## 2.2、jmeter\_slaves

jmeter\_slaves\_deploy.yaml (Jmeter slave 的部署清单) :

```
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: jmeter-slaves
5    labels:
6      jmeter_mode: slave
7  spec:
8    replicas: 2
9    selector:
10     matchLabels:
11       jmeter_mode: slave
12    template:
13     metadata:
14       labels:
15         jmeter_mode: slave
16     spec:
17       containers:
18       - name: jmslave
19         image: registry.cn-beijing.aliyuncs.com/7d/jmeter-slave:latest
20         imagePullPolicy: IfNotPresent
21         ports:
22         - containerPort: 1099
23         - containerPort: 50000
24       resources:
25         limits:
26           cpu: 4000m
27           memory: 4Gi
28         requests:
29           cpu: 500m
30           memory: 512Mi
```

jmeter\_slaves\_svc.yaml ( Jmeter slave 的服务清单) :

```
1  apiVersion: v1
2  kind: Service
3  metadata:
4    name: jmeter-slaves-svc
5    labels:
6      jmeter_mode: slave
7  spec:
8    clusterIP: None
9    ports:
10     - port: 1099
11       name: first
12       targetPort: 1099
13     - port: 50000
14       name: second
15       targetPort: 50000
```

## 2.3、jmeter\_master

jmeter\_master\_configmap.yaml (jmeter\_master 应用配置) :

```
1 apiVersion: v1
2 kind: ConfigMap
3 metadata:
4   name: jmeter-load-test
5   labels:
6     app: influxdb-jmeter
7 data:
8   load_test: |
9     #!/bin/bash
10    #Script created to invoke jmeter test script with the slave POD IP addresses
11    #Script should be run like: ./load_test "path to the test script in jmx format"
12    /jmeter/apache-jmeter-*/bin/jmeter -n -t $1 `getent ahostsv4 jmeter-slaves-svc | cut -d' ' -f1 | sort
```

jmeter\_master\_deploy.yaml (jmeter\_master 部署清单) :

```
1 apiVersion: apps/v1 # for versions before 1.9.0 use apps/v1beta2
2 kind: Deployment
3 metadata:
4   name: jmeter-master
5   labels:
6     jmeter_mode: master
7 spec:
8   replicas: 1
9   selector:
10    matchLabels:
11      jmeter_mode: master
12   template:
13     metadata:
14       labels:
15         jmeter_mode: master
16     spec:
17       containers:
18         - name: jmmaster
19           image: registry.cn-beijing.aliyuncs.com/7d/jmeter-master:latest
20           imagePullPolicy: IfNotPresent
21           command: [ "/bin/bash", "-c", "--" ]
22           args: [ "while true; do sleep 30; done;" ]
23           volumeMounts:
24             - name: loadtest
25               mountPath: /load_test
26               subPath: "load_test"
27           ports:
28             - containerPort: 60000
29           resources:
30             limits:
31               cpu: 4000m
32               memory: 4Gi
33             requests:
34               cpu: 500m
35               memory: 512Mi
36       volumes:
37         - name: loadtest
38           configMap:
39             name: jmeter-load-test
```

## 2.4、influxdb

jmeter\_influxdb\_configmap.yaml (influxdb 的应用配置) :

```
1 apiVersion: v1
2 kind: ConfigMap
3 metadata:
4   name: influxdb-config
5   labels:
6     app: influxdb-jmeter
7 data:
8   influxdb.conf: |
9     [meta]
10    dir = "/var/lib/influxdb/meta"
11
12    [data]
13    dir = "/var/lib/influxdb/data"
14    engine = "tsml"
15    wal-dir = "/var/lib/influxdb/wal"
16
17    # Configure the graphite api
18    [[graphite]]
19    enabled = true
20    bind-address = ":2003" # If not set, is actually set to bind-address.
21    database = "jmeter" # store graphite data in this database
```

jmeter\_influxdb\_deploy.yaml (influxdb 部署清单) :

```
1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   name: influxdb-jmeter
5   labels:
6     app: influxdb-jmeter
7 spec:
8   replicas: 1
9   selector:
10    matchLabels:
11      app: influxdb-jmeter
12   template:
13     metadata:
14       labels:
15         app: influxdb-jmeter
16     spec:
17       containers:
18         - image: influxdb
19           imagePullPolicy: IfNotPresent
20           name: influxdb
21           volumeMounts:
22             - name: config-volume
23               mountPath: /etc/influxdb
24           ports:
25             - containerPort: 8083
26               name: influx
27             - containerPort: 8086
28               name: api
29             - containerPort: 2003
30               name: graphite
31       volumes:
32         - name: config-volume
33           configMap:
34             name: influxdb-config
```



jmeter\_influxdb\_svc.yaml (influxdb 部署服务清单) :

```
1  apiVersion: v1
2  kind: Service
3  metadata:
4    name: jmeter-influxdb
5    labels:
6      app: influxdb-jmeter
7  spec:
8    ports:
9      - port: 8083
10       name: http
11       targetPort: 8083
12      - port: 8086
13       name: api
14       targetPort: 8086
15      - port: 2003
16       name: graphite
17       targetPort: 2003
18  selector:
19    app: influxdb-jmeter
```

## 2.5、grafana

jmeter\_grafana\_deploy.yaml (grafana 部署清单) :

```
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: jmeter-grafana
5    labels:
6      app: jmeter-grafana
7  spec:
8    replicas: 1
9    selector:
10     matchLabels:
11       app: jmeter-grafana
12    template:
13     metadata:
14       labels:
15         app: jmeter-grafana
16     spec:
17       containers:
18         - name: grafana
19           image: grafana/grafana:5.2.0
20           imagePullPolicy: IfNotPresent
21           ports:
22             - containerPort: 3000
23               protocol: TCP
24           env:
25             - name: GF_AUTH_BASIC_ENABLED
26               value: "true"
27             - name: GF_USERS_ALLOW_ORG_CREATE
28               value: "true"
29             - name: GF_AUTH_ANONYMOUS_ENABLED
30               value: "true"
31             - name: GF_AUTH_ANONYMOUS_ORG_ROLE
32               value: Admin
33             - name: GF_SERVER_ROOT_URL
34               # If you're only using the API Server proxy, set this value instead:
35               # value: /api/v1/namespaces/kube-system/services/monitoring-grafana/proxy
36               value: /
```

jmeter\_grafana\_svc.yaml (grafana 部署服务清单) :

```
1  apiVersion: v1
2  kind: Service
3  metadata:
4    name: jmeter-grafana
5    labels:
6      app: jmeter-grafana
7  spec:
8    ports:
9      - port: 3000
10      targetPort: 3000
11    selector:
12      app: jmeter-grafana
13    type: NodePort
14  ---
15  apiVersion: extensions/v1beta1
16  kind: Ingress
17  metadata:
18    annotations:
19      nginx.ingress.kubernetes.io/service-weight: 'jmeter-grafana: 100'
20    name: jmeter-grafana-ingress
21  spec:
22    rules:
23      # 配置七层域名
24      - host: grafana-jmeter.7d.com
25        http:
26          paths:
27            # 配置Context Path
28            - path: /
29              backend:
30                serviceName: jmeter-grafana
31                servicePort: 3000
```

## 五、初始化 dashboard

### 1、启动 dashboard 脚本

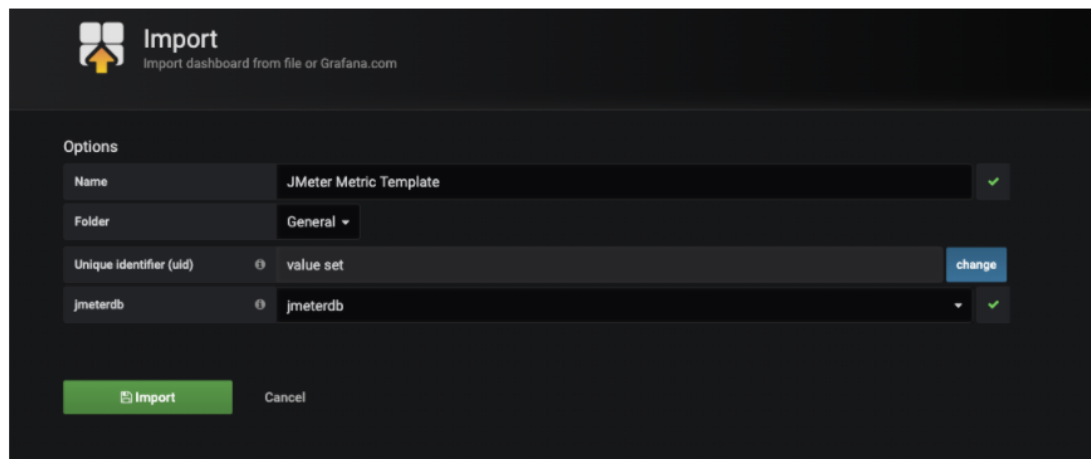
```
1 | $ ./dashboard.sh
```

检查 service 部署情况:

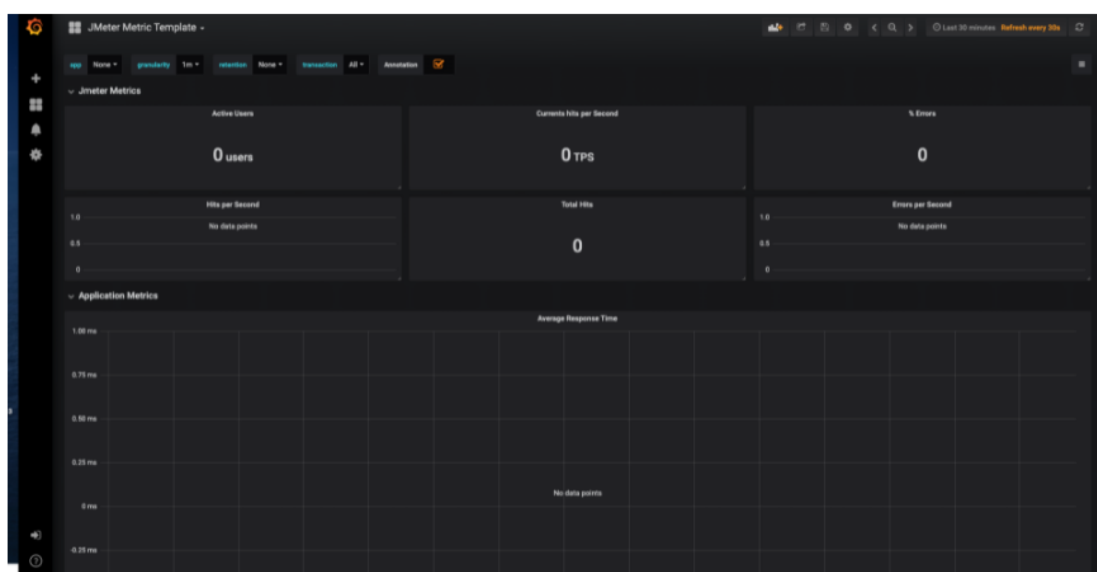
```
1 | $ kubectl get svc -n 7dgroup
2 | NAME                                TYPE                CLUSTER-IP      EXTERNAL-IP      PORT(S)          AGE
3 | jmeter-grafana                      NodePort            10.96.6.201     <none>           3000:31801/TCP   10m
4 | jmeter-influxdb                    ClusterIP           10.96.111.60    <none>           8083/TCP,8086/TCP,2003/TCP 10m
5 | jmeter-slaves-svc                  ClusterIP           None            <none>           1099/TCP,50000/TCP 10m
```

我们可以通过 [http://任意\\_node\\_ip:31801/](http://任意_node_ip:31801/) 访问 grafana

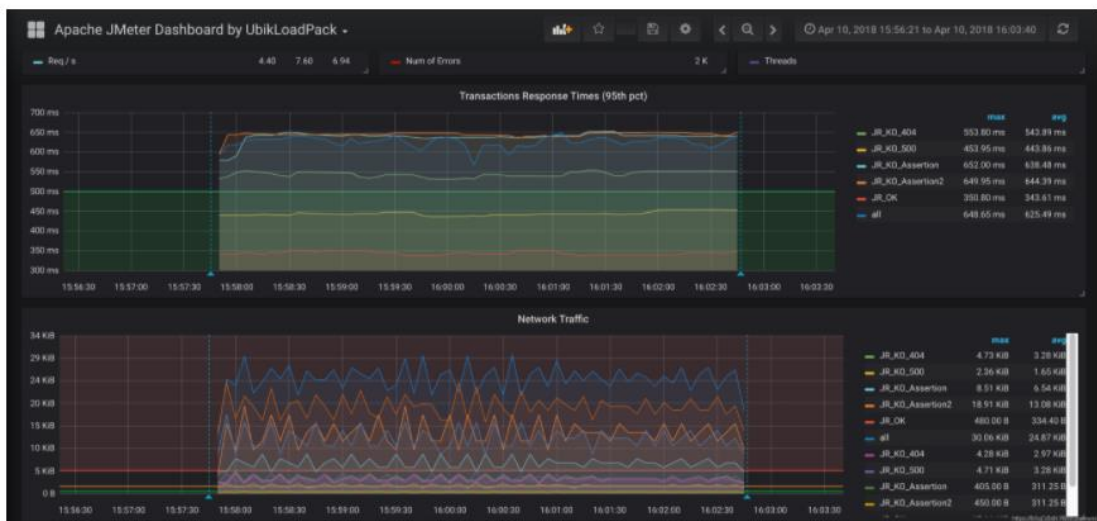
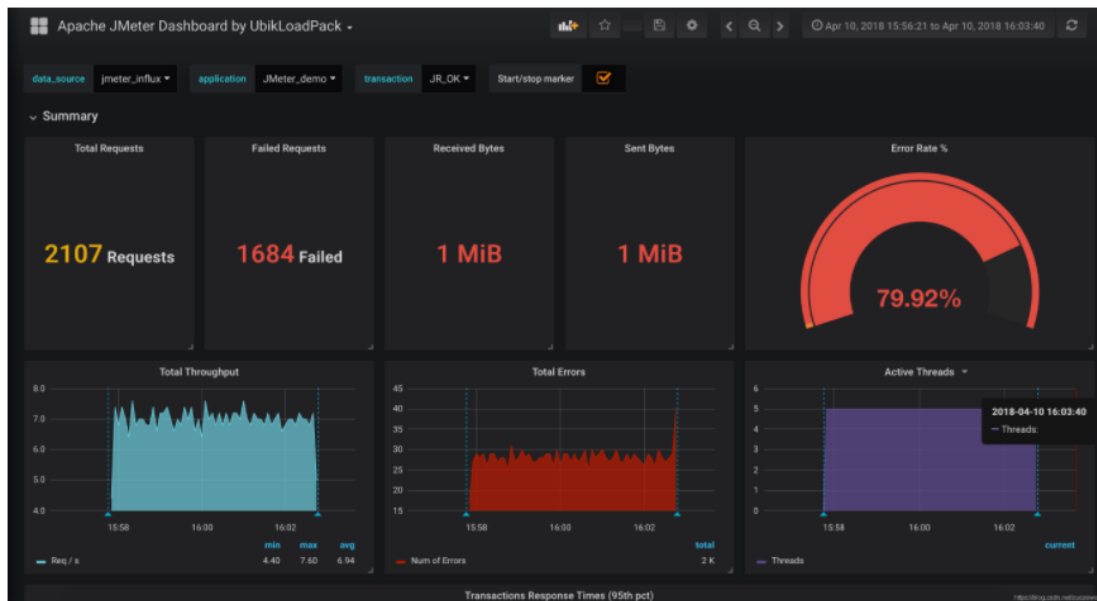
最后, 我们在 grafana 导入 dashborad 模版:



The image shows the Grafana 'Import' dashboard screen. At the top, there's a header with the Grafana logo and the text 'Import dashboard from file or Grafana.com'. Below this, there's a section titled 'Options' with several input fields: 'Name' (JMeter Metric Template), 'Folder' (General), 'Unique identifier (uid)' (value set), and 'jmeterdb' (jmeterdb). There are green checkmarks next to the 'Name' and 'jmeterdb' fields. At the bottom, there are two buttons: 'Import' and 'Cancel'.



如果你不喜欢这个模版, 也可以导入热门模版: [5496](#)



## 2、部署清单

`dashboard.sh` 该脚本用于自动创建以下内容:

- (1) influxdb pod 中的一个 influxdb 数据库 (Jmeter)
- (2) grafana 中的数据源 (jmeterdb)

```
1  #!/usr/bin/env bash
2  working_dir=`pwd`
3  #Get namespace variable
4  tenant=`awk '{print $NF}' $working_dir/tenant_export`
5  ## Create jmeter database automatically in Influxdb
6  echo "Creating Influxdb jmeter Database"
7  ##Wait until Influxdb Deployment is up and running
8  ##influxdb_status=`kubectl get po -n $tenant | grep influxdb-jmeter | awk '{print $2}' | grep Running`
9  influxdb_pod=`kubectl get po -n $tenant | grep influxdb-jmeter | awk '{print $1}'`
10 kubectl exec -ti -n $tenant $influxdb_pod -- influx -execute 'CREATE DATABASE jmeter'
11 ## Create the influxdb datasource in Grafana
12 echo "Creating the Influxdb data source"
13 grafana_pod=`kubectl get po -n $tenant | grep jmeter-grafana | awk '{print $1}'`
14 ## Make load test script in Jmeter master pod executable
15 #Get Master pod details
16 master_pod=`kubectl get po -n $tenant | grep jmeter-master | awk '{print $1}'`
17 kubectl exec -ti -n $tenant $master_pod -- cp -r /load_test /[]()jmeter/load_test
18 kubectl exec -ti -n $tenant $master_pod -- chmod 755 /jmeter/load_test
19 ##kubectl cp $working_dir/influxdb-jmeter-datasource.json -n $tenant $grafana_pod:/influxdb-jmeter-datasou
20 kubectl exec -ti -n $tenant $grafana_pod -- curl 'http://admin:admin@127.0.0.1:3000/api/datasources' -X PC
```

## 六、启动测试

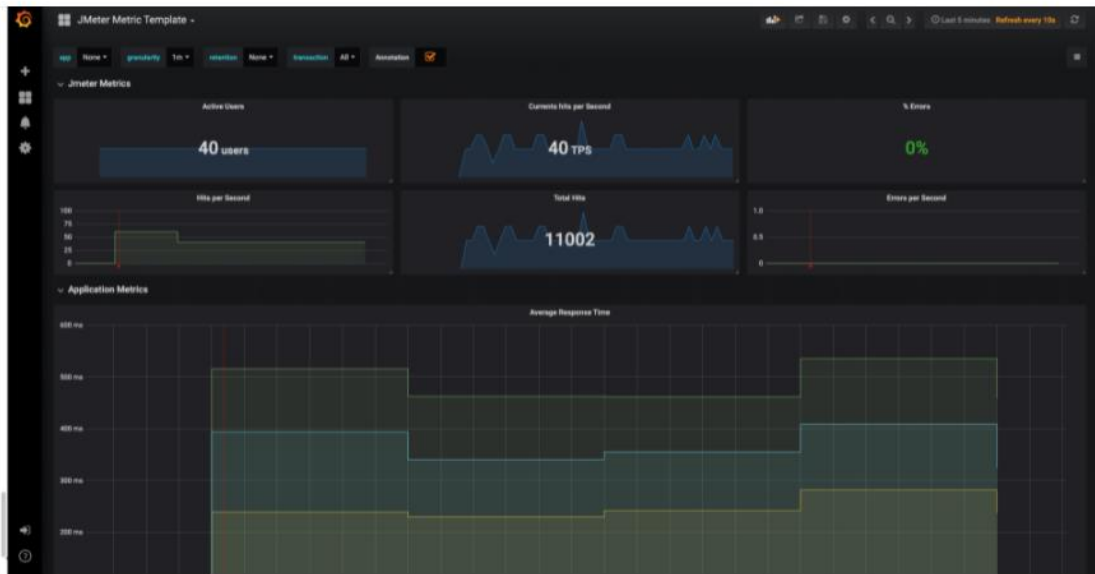
### 1、执行脚本

```
1 | $ ./start_test.sh
```

需要一个测试脚本，本例为： `web-test.jmx`

```
1 | $ ./start_test.sh
2 | Enter path to the jmx file web-test.jmx
3 | 'SLF4J: Class path contains multiple SLF4J bindings.
4 | SLF4J: Found binding in [jar:file:/jmeter/apache-jmeter-5.0/lib/log4j-slf4j-impl-2.11.0.jar!/org/slf4j/imp
5 | SLF4J: Found binding in [jar:file:/jmeter/apache-jmeter-5.0/lib/ext/pepper-box-1.0.jar!/org/slf4j/impl/Sta
6 | SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
7 | SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
8 | Jul 25, 2020 11:30:58 AM java.util.prefs.FileSystemPreferences$1 run
9 | INFO: Created user preferences directory.
10 | Creating summariser <summary>
11 | Created the tree successfully using web-test.jmx
12 | Configuring remote engine: 10.100.113.31
13 | Configuring remote engine: 10.100.167.173
14 | Starting remote engines
15 | Starting the test @ Sat Jul 25 11:30:59 UTC 2020 (1595676659540)
16 | Remote engines have been started
17 | Waiting for possible Shutdown/StopTestNow/Heapdump message on port 4445
18 | summary + 803 in 00:00:29 = 27.5/s Avg: 350 Min: 172 Max: 1477 Err: 0 (0.00%) Active: 40 Sta
19 | summary + 1300 in 00:00:29 = 45.3/s Avg: 367 Min: 172 Max: 2729 Err: 0 (0.00%) Active: 40 Sta
20 | summary = 2103 in 00:00:58 = 36.4/s Avg: 361 Min: 172 Max: 2729 Err: 0 (0.00%)
21 | summary + 1400 in 00:00:31 = 45.4/s Avg: 342 Min: 160 Max: 2145 Err: 0 (0.00%) Active: 40 Sta
22 | summary = 3503 in 00:01:29 = 39.5/s Avg: 353 Min: 160 Max: 2729 Err: 0 (0.00%)
23 | summary + 1400 in 00:00:31 = 45.2/s Avg: 352 Min: 169 Max: 2398 Err: 0 (0.00%) Active: 40 Sta
24 | summary = 4903 in 00:02:00 = 41.0/s Avg: 353 Min: 160 Max: 2729 Err: 0 (0.00%)
25 | summary + 1400 in 00:00:30 = 46.8/s Avg: 344 Min: 151 Max: 1475 Err: 0 (0.00%) Active: 40 Sta
26 | summary = 6303 in 00:02:30 = 42.1/s Avg: 351 Min: 151 Max: 2729 Err: 0 (0.00%)
27 | summary + 1200 in 00:00:28 = 43.5/s Avg: 354 Min: 163 Max: 2018 Err: 0 (0.00%) Active: 40 Sta
28 | summary = 7503 in 00:02:57 = 42.3/s Avg: 351 Min: 151 Max: 2729 Err: 0 (0.00%)
29 | summary + 1300 in 00:00:30 = 43.7/s Avg: 456 Min: 173 Max: 2401 Err: 0 (0.00%) Active: 40 Sta
30 | summary = 8803 in 00:03:27 = 42.5/s Avg: 367 Min: 151 Max: 2729 Err: 0 (0.00%)
31 | summary + 1400 in 00:00:31 = 44.9/s Avg: 349 Min: 158 Max: 2128 Err: 0 (0.00%) Active: 40 Sta
32 | summary = 10203 in 00:03:58 = 42.8/s Avg: 364 Min: 151 Max: 2729 Err: 0 (0.00%)
33 | summary + 1400 in 00:00:32 = 44.3/s Avg: 351 Min: 166 Max: 1494 Err: 0 (0.00%) Active: 40 Sta
```

查看测试数据：



## 2、部署清单

`start_test.sh`（此脚本用于自动运行 Jmeter 测试脚本，而无需手动登录 Jmeter 主 shell，它将询问 Jmeter 测试脚本的位置，然后将其复制到 Jmeter master pod 并启动自动对 Jmeter slave 进行测试）：

```
1  #!/usr/bin/env bash
2  #Script created to launch Jmeter tests directly from the current terminal without accessing the jmeter mas
3  #It requires that you supply the path to the jmx file
4  #After execution, test script jmx file may be deleted from the pod itself but not locally.
5
6  #直接从当前终端启动 Jmeter 测试而创建的脚本，无需访问 Jmeter master pod。
7  #要求提供 jmx 文件的路径
8  #执行后，测试脚本 jmx 文件可能会从 pod 本身删除，但不会在本地删除。
9
10 working_dir="`pwd`"
11
12 # 获取 namespace 变量
13 tenant=`awk '{print $NF}' "$working_dir/tenant_export"`
14
15 jmx="$1"
16 [ -n "$jmx" ] || read -p 'Enter path to the jmx file ' jmx
17
18 if [ ! -f "$jmx" ];
19 then
20     echo "Test script file was not found in PATH"
21     echo "Kindly check and input the correct file path"
22     exit
23 fi
24
25 test_name="$(basename "$jmx")"
26
27 # 获取 master pod 详细信息
28 master_pod=`kubectl get po -n $tenant | grep jmeter-master | awk '{print $1}'`
29 kubectl cp "$jmx" -n $tenant "$master_pod:$test_name"
30
31 ## 启动 Jmeter 压测
32 kubectl exec -ti -n $tenant $master_pod -- /bin/bash /load_test "$test_name"
33 kubectl exec -ti -n $tenant $master_pod -- /bin/bash /load_test "$test_name"
```

jmeter\_stop.sh (停止测试) :

```
1  #!/usr/bin/env bash
2  #Script writtent to stop a running jmeter master test
3  #Kindly ensure you have the necessary kubeconfig
4
5  #编写脚本来停止运行的 jmeter master 测试
6  #请确保你有必要的 kubeconfig
7  working_dir=`pwd`
8
9  #获取 namespace 变量
10 tenant=`awk '{print $NF}' $working_dir/tenant_export`
11 master_pod=`kubectl get po -n $tenant | grep jmeter-master | awk '{print $1}'`
12 kubectl -n $tenant exec -it $master_pod -- bash -c "./jmeter/apache-jmeter-5.0/bin/stoptest.sh"
```

## 七、小结

传统 Jmeter 存在的问题：

- 并发数超过单节点承载能力时，多节点环境配置、维护复杂；
- 默认配置下无法并行运行多个测试，需要更改配置启动额外进程；
- 难以支持云环境下测试资源的弹性伸缩需求。

Kubernetes-Jmeter 带来的改变：

- 压测执行节点一键安装；
- 多个项目、多个测试可并行使用同一个测试资源池（最大并发数允许情况下，Kubernetes 也提供了 RBAC、namespace 等管理能力，支持多用户共享一个集群，并实现资源限制），提高资源利用率；
- 对接 Kubernetes HPA 根据并发数自动启动、释放压测执行节点。

源码地址：

- <https://github.com/zuozewei/blog-example/tree/master/Kubernetes/k8s-jmeter-cluster>

参考资料：

- [1]: <https://github.com/kubernauts/jmeter-kubernetes>